

What an Apify Starter plan actually buys the lead-gen agent

A breakdown of current SDR pipeline costs — lead sourcing, email verification, and AI token usage — plus a sizing model for how many companies and leads a \$29/month Apify Starter plan can realistically support.

Prepared for	Date	System
Management review	August 20, 2026	AI SDR Agent (ai_sdr_agent)

Status
Apify is already integrated — this sizes usage, not a new build

Executive summary

The SDR agent already has a working Apify integration — this is not a new build. On the **Starter plan (\$29/mo)**, using the lead-sourcing actor already wired into the system that returns verified emails, the account can pull roughly **19,000 leads with emails per month** before the credit runs out — about **640/day**. That's the ceiling on the data-sourcing side. The real bottleneck is elsewhere: the sending pipeline caps every lead at **one email per day system-wide**, and AI qualification currently processes leads in batches of 5 every 2 minutes — about **150/hour** — so a burst of 19,000 fresh leads would take roughly **5.3 days** just to get qualified, long before it becomes an outreach bottleneck.

LEADS WITH EMAIL / MONTH

~19,300

on \$29 Starter credit, \$1.50/1k actor

COMPANIES COVERABLE / MONTH

~3,800–6,400

at 3–5 leads sourced per company

COST PER LEAD (SOURCING)

\$0.0015

Apify actor cost only

QUALIFICATION THROUGHPUT

150/hr

shared across all companies on the platform

01 Current pipeline — where cost actually enters

Confirmed from the codebase

(`ai_sdr_agent/agents/lead_research_agent.py`, `tasks.py`, `outreach_agent.py`). There is no separate "companies" table — company fields live directly on each lead row.

1

Source leads

Priority order: Apollo.io API → **Apify actor** → Groq AI-generated fallback. Apify path uses `code_crafter/leads-finder` (\$1.50/1k, returns verified email) or the default `apify/google-search-scraper` (LinkedIn search results only, **no email field**).

2

Format-validate

Regex/heuristic check only (`lead_validator.py`) — flags fake or free-domain addresses, scores 0–100. Not a real deliverability check. \$0 cost, runs locally.

3

AI qualification

Groq LLM scores each lead (hot/warm/cold). Batches of **5 leads every 120 seconds**, shared queue across the whole platform — not per company. Token usage metered.

4

Enroll in campaign

Manual selection → `SDRCampaignEnrollment` . No cap found on leads per campaign or per enroll call.

5

Send & track

Per-campaign SMTP. Hard guard: **max 1 email per lead per day**, system-wide across all campaigns. Bounces detected after send via IMAP reply parsing — not pre-send verification.

⚠ No email verification service is integrated today

The system checks email *format* (regex) and detects *bounces* after a send fails — but there is no Hunter.io, NeverBounce, or ZeroBounce-style pre-send verification anywhere in the codebase. The `leads-finder`

Apify actor's own "verified" email field is the closest thing to it currently, and it's actor-side, not ours. Adding real pre-send verification (see §4) would be a new integration.

02 Apify Starter plan — capacity model

Verified current pricing directly from apify.com (the in-app setup guide still shows an outdated "\$49/month" — actual current price is below).

Apify plan tiers – apify.com/pricing, verified August 2026					
PLAN	MONTHLY PRICE	USAGE CREDIT	CONCURRENT RUNS	NOTES	
Free	\$0	\$5	low	Can only import from existing finished r	
Starter	\$29	\$29	32	Credit does not roll over — resets mont	
Scale	\$199	\$199	higher	For reference only — not requested	

Actor cost — the number that actually decides lead volume

The plan's dollar amount only matters relative to the actor's per-result price. Two actors are already referenced in this codebase; only one produces usable emails:

code_crafter/leads-finder		apify/google-search-scraper	
\$1.50 per 1,000 leads · pay-per-event		~\$0.15–0.50 per 1,000 results · default actor	
Returns verified email	Yes	Returns verified email	No
Leads on \$29 credit	~19,333	Results on \$29 credit	~58,000–193,000
Per day (30-day month)	~644	Per day (30-day month)	~1,930–6,400
Per minute (continuous)	~0.45	Usable as-is for outreach	No — name/title/company only
Cost per lead	\$0.0015		

Reading "how many leads per minute"

0.45 leads/minute is the *credit-exhaustion* rate if Apify ran continuously all month — not how the system actually behaves. In practice Apify runs happen in batches (current default: 10 leads per scheduled run, one run every 24 hours via the background auto-research job). A single manual run against leads-finder can return hundreds of leads in under two minutes — the \$29 credit, not per-minute throughput, is what limits monthly volume.

Companies coverable per month

"Companies" isn't a tracked entity in this system — each lead carries its own company fields. Coverage depends on how many contacts you pull per target company:

Same \$29 credit, same 19,333 leads – sliced by contacts-per-company

LEADS PULLED PER COMPANY	COMPANIES COVERED / MONTH	COMPANIES COVERED / DAY
1 (single decision-maker)	19,333	644
3 (typical: 2–3 contacts)	6,444	215
5 (broader account mapping)	3,867	129

03 Downstream bottlenecks that cap the Apify number

Sourcing 19,000 leads a month doesn't mean 19,000 leads get emailed — two hard limits already in the code throttle everything after sourcing.

Rate limits found in ai_sdr_agent/tasks.py and scheduler.py

STAGE	LIMIT	EFFECTIVE THROUGHPUT
AI qualification (Groq)	5 leads / 120s, one shared queue	~150/hr · ~3,600/day
Email sending	1 email per lead per day, all campaigns combined	1/lead/day max
Send scheduler tick	runs every 5 minutes	288 ticks/day
Auto lead research	10 leads / run, 1 run / 24h	10/day (auto path only)

⚠ Auto-research scheduling gap

Two separate auto-research triggers exist in code — a background thread (`middleware.py`, every 24h) and a Celery task (`tasks.py`) — but the active APScheduler job list does not register the Celery version. Only the thread-based path appears to run today. Worth a dev ticket regardless of the Apify decision.

Net effect: qualification throughput (~3,600/day max) comfortably exceeds Apify's own daily ceiling (~644/day at leads-finder pricing), so **Apify credit — not internal processing — is the binding constraint** on lead volume today.

04 Email verification cost — if added

No verification provider is integrated today (\$1). If the manager wants real pre-send deliverability checks (catch-all detection, SMTP ping, bounce prediction) rather than relying on the Apify actor's own "verified" flag, Hunter.io is the common reference point:

Hunter.io — verification-only credits (0.5 credit per email checked)

PLAN	MONTHLY PRICE	CREDITS	EMAILS VERIFIABLE	COST / VERIFIED EMAIL
Starter	~\$34/mo (annual)	2,000	4,000	~\$0.0085
Growth	~\$104/mo (annual)	10,000	20,000	~\$0.0052
Scale	~\$209/mo (annual)	25,000	50,000	~\$0.0042

Adding verification on top of the leads-finder actor's ~19,000 monthly leads would need a plan covering ~19,000 checks — the **Growth tier (~\$104/mo)** comfortably covers that with headroom, bringing blended sourcing + verification cost to roughly **\$0.007/lead** total.

05 AI usage cost — what's tracked today

The platform already meters LLM token usage per company per agent (core/api_key_service.py), but it is **not broken out by campaign or by lead** anywhere in the system. Every SDR sub-agent that calls the model — qualification, outreach personalization, reply/meeting handling — records into the same company-level pool.

FREE TOKENS / COMPANY / AGENT

10,000

DEFAULT_FREE_TOKENS, granted on purchase

MANAGED-KEY RESET CYCLE

30 days

as configured for this key (reset_interval_days)

PER-CAMPAIGN COST TRACKING

None

not present in current models

PER-LEAD COST TRACKING

None

not present in current models

APIFY SPEND VISIBILITY

None

billed entirely on Apify's own dashboard

Gap for the manager to weigh

Right now, nobody inside the app can see "campaign X cost \$Y in AI + \$Z in Apify credit." Token spend is visible per company, but Apify spend is invisible until you check Apify's own billing dashboard, and neither is attributed to a specific campaign or lead. If per-campaign ROI reporting matters, that's a small but real gap to close alongside the Apify plan decision.

06 **Bottom line**

All-in monthly cost, Starter plan + leads-finder actor + optional verification		
COMPONENT	MONTHLY COST	COVERS
Apify Starter plan	\$29	~19,300 leads with actor
Hunter.io Growth (optional, if adding real verification)	~\$104	20,000 verification check
AI token usage (SDR agents)	\$0 up to free tier	10,000 free tokens/comp

COMPONENT	MONTHLY COST	COVERS
Total (sourcing only)	\$29/mo	~644 leads/day, ~215 cc
Total (sourcing + verification)	~\$133/mo	same volume, with indep

Recommendation: start with the **\$29 Starter plan alone** — the leads-finder actor's built-in "verified" email flag already covers the basic case, and current qualification/send throughput has plenty of headroom above Apify's ~644 leads/day ceiling. Add Hunter.io verification only if bounce rates from actor-side verification prove too high in practice.



07 Billing scenario — Apify credit + real AI enrichment cost

AI cost only applies to the **enrichment step** — the qualification agent that scores each lead (`lead_qualification_agent.py`) — not to sourcing or verification. That call runs on Groq's `llama-3.1-8b-instant` (confirmed in code, no override configured), priced at \$0.05/1M input and \$0.08/1M output tokens. One qualification call is ~750 input + ~300 output tokens, so real cost is about **\$0.00006 per lead enriched** — not a flat \$1. Volume is fixed by the \$29 Apify credit either way.

Leads/month fixed by the \$29 Apify credit (\$1.50/1k, leads-finder actor) · AI enrichment at actual Groq cost, ~\$0.00006/lead			
PLAN COMBINATION	LEADS / MONTH	AI ENRICHMENT COST	TOTAL MONTHLY COST
Apify Starter only \$29/mo, no verification add-on	19,333	~\$1.16	~\$30.16
Apify Starter + Hunter.io Growth \$29 + ~\$104/mo verification	19,333	~\$1.16	~\$134.16

AI enrichment cost is negligible next to Apify/Hunter.io

At Groq's real per-token rate, enriching all 19,333 leads costs roughly **\$1.16 for the whole month** — under 4% of the \$29 Apify plan itself, and it's covered almost entirely by the platform's free token tier (10,000 free tokens/company) before any paid usage kicks in. The Apify credit — not AI processing — is what actually

caps monthly volume and cost. Outreach personalization (a separate Groq call per campaign step actually sent, only when `ai_personalize` is on) and reply-classification calls add a small amount on top, but at the same sub-cent-per-call rate.

Sources: apify.com/pricing (fetched Aug 2026) · apify.com/code_crafter/leads-finder · Hunter.io published pricing (2026) · internal codebase review of `ai_sdr_agent/` (`models.py`, `tasks.py`, `lead_research_agent.py`, `outreach_agent.py`, `scheduler.py`, `api_key_service.py`)